

# Retail Alpha Forecaster - Notebook 1: Exploratory Data Analysis

## Notebook 1 — EDA (Exploratory Data Analysis)

In this notebook, we explored the raw retail sales data to understand trends, seasonality, and sparsity. We identified that demand is highly intermittent with many zero-sales days, which affects error metrics like MAPE/SMAPE. Visual analysis confirmed strong weekly and monthly seasonality. This foundation guided feature engineering for the forecasting pipeline.

```
In [5]: # --- Google Cloud / BigQuery Setup ---
# --- Cell 1: imports & config ---
from pathlib import Path
import os, sys, math
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from google.cloud import bigquery
from google.oauth2 import service_account

print("Python:", sys.executable)
print("CWD      :", Path.cwd())

# ---- Project / dataset / tables ----
PROJECT = "retail-alpha-forecaster"
DATASET = "raf" # keep consistent across notebooks
RAW_TABLE = f"`{PROJECT}.{DATASET}.raw_sales`"

# ---- Service-account JSON resolution (works both in VS Code & browser) ----
KEY_FILENAME = "retail-alpha-forecaster-7f14a7b50e62.json"
CANDIDATES = [
    Path.cwd() / "keys" / KEY_FILENAME, # running repo root
    Path.cwd().parents[0] / "keys" / KEY_FILENAME, # running inside notebook
    Path.cwd().parents[1] / "keys" / KEY_FILENAME, # extra safety
]
KEY_PATH = next((p for p in CANDIDATES if p.exists()), None)
assert KEY_PATH and KEY_PATH.exists(), f"Key not found. Looked for: {KEY_FILENAME}"

# either let google libs pick up env var...
os.environ["GOOGLE_APPLICATION_CREDENTIALS"] = str(KEY_PATH)
client = bigquery.Client(project=PROJECT)

# or build creds explicitly (both ways are fine)
# creds = service_account.Credentials.from_service_account_file(str(KEY_PATH))
# client = bigquery.Client(project=PROJECT, credentials=creds)

def q(sql: str) -> pd.DataFrame:
```

```
"""Run a BQ SQL string and return a pandas DataFrame."""  
return client.query(sql).result().to_dataframe()
```

Python: /home/btheard/retail-alpha-forecaster/.venv/bin/python

CWD : /home/btheard/retail-alpha-forecaster/notebooks

In [6]: # --- Cell 2: load typed & filtered data ---

```
USE_SAMPLE = False      # <- flip to True for quick dev  
SAMPLE_N    = 5000      # ignored if USE_SAMPLE=False  
  
# Window we decided for this project (the Kaggle time range)  
DATE_START = "2013-01-01"  
DATE_END   = "2015-10-31"  
  
LIMIT_CLAUSE = "" if not USE_SAMPLE else f"LIMIT {SAMPLE_N}"  
  
sql = f"""  
WITH base AS (  
    SELECT  
        SAFE_CAST(date AS DATE)           AS date,  
        SAFE_CAST(date_block_num AS INT64) AS date_block_num,  
        SAFE_CAST(shop_id AS INT64)       AS shop_id,  
        SAFE_CAST(item_id AS INT64)       AS item_id,  
        SAFE_CAST(item_price AS FLOAT64)  AS item_price,  
        SAFE_CAST(item_cnt_day AS FLOAT64) AS item_cnt_day  
    FROM {RAW_TABLE}  
)  
clean AS (  
    SELECT  
        date, date_block_num, shop_id, item_id, item_price, item_cnt_day,  
        -- tag returns (item_cnt_day < 0) without deleting them  
        (item_cnt_day < 0) AS is_return  
    FROM base  
    WHERE shop_id IS NOT NULL  
        AND item_id IS NOT NULL  
        AND date BETWEEN DATE('{DATE_START}') AND DATE('{DATE_END}'))  
SELECT * FROM clean  
ORDER BY date, shop_id, item_id  
{LIMIT_CLAUSE}  
"""  
  
df = q(sql)  
df["date"] = pd.to_datetime(df["date"])  
print(f"Rows: {len(df):,}")  
df.info()  
df.head()
```

```

Rows: 2,935,849
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2935849 entries, 0 to 2935848
Data columns (total 7 columns):
#   Column                Dtype
---  -
0   date                  datetime64[ns]
1   date_block_num        Int64
2   shop_id               Int64
3   item_id               Int64
4   item_price            float64
5   item_cnt_day          float64
6   is_return             boolean
dtypes: Int64(3), boolean(1), datetime64[ns](1), float64(2)
memory usage: 148.4 MB

```

```

Out[6]:

```

	date	date_block_num	shop_id	item_id	item_price	item_cnt_day	is_return
0	2013-01-01	0	2	991	99.0	1.0	False
1	2013-01-01	0	2	1472	2599.0	1.0	False
2	2013-01-01	0	2	1905	249.0	1.0	False
3	2013-01-01	0	2	2920	599.0	2.0	False
4	2013-01-01	0	2	3320	1999.0	1.0	False

```

In [7]: # --- Cell 3: coverage & sanity ---

print("Unique shops :", df["shop_id"].nunique())
print("Unique items :", df["item_id"].nunique())

dmin, dmax = df["date"].min(), df["date"].max()
print("Date range :", dmin.date(), "→", dmax.date())

# Expect exactly 34 months (2013-01 .. 2015-10). Confirm using date_block_num
months_present = sorted(df["date_block_num"].unique())
print("Distinct months (#):", len(months_present))
missing_blocks = set(range(min(months_present), max(months_present)+1)) - set(months_present)
print("Missing month blocks:", missing_blocks)

# Nulls?
nulls = df.isna().sum().sort_values(ascending=False)
display(nulls[nulls > 0])

# Negative prices? zero prices? (pay attention for data quirks)
n_neg_price = (df["item_price"] < 0).sum()
n_zero_price = (df["item_price"] == 0).sum()
print(f"Negative price rows: {n_neg_price:,} | Zero price rows: {n_zero_price:,}")

# Returns (item_cnt_day < 0) – tag from SQL; decide handling later
n_returns = df["is_return"].sum()

```

```

print(f"Return rows (tagged): {n_returns:,}")

# Duplicates for (date, shop_id, item_id) (shouldn't happen; forecast target
dups = (df.groupby(["date", "shop_id", "item_id"])
        .size().reset_index(name="cnt")
        .query("cnt > 1")
        .sort_values("cnt", ascending=False))
print("Duplicate day-shop-item keys:", len(dups))
dups.head(10)

```

Unique shops : 60  
 Unique items : 21807  
 Date range : 2013-01-01 → 2015-10-31  
 Distinct months (#): 34  
 Missing month blocks: set()  
 Series([], dtype: int64)  
 Negative price rows: 1 | Zero price rows: 0  
 Return rows (tagged): 7,356  
 Duplicate day-shop-item keys: 28

Out[7]:

	date	shop_id	item_id	cnt
<b>25568</b>	2013-01-05	54	20130	2
<b>94095</b>	2013-01-25	31	14050	2
<b>2460279</b>	2015-02-21	25	16587	2
<b>2449167</b>	2015-02-17	5	21619	2
<b>2436443</b>	2015-02-12	42	21619	2
<b>2322738</b>	2014-12-31	57	8237	2
<b>2320472</b>	2014-12-31	42	21619	2
<b>2279362</b>	2014-12-26	17	3424	2
<b>2264631</b>	2014-12-22	31	8237	2
<b>2163951</b>	2014-11-21	31	16587	2

In [8]:

```

# --- Cell 4: basic distributions & sparsity ---

sns.set_theme(style="whitegrid")

# Price distribution (log-x often more readable)
plt.figure(figsize=(8,4))
sns.histplot(df["item_price"], bins=100, kde=False)
plt.title("Item price distribution")
plt.xlabel("price"); plt.ylabel("rows")
plt.show()

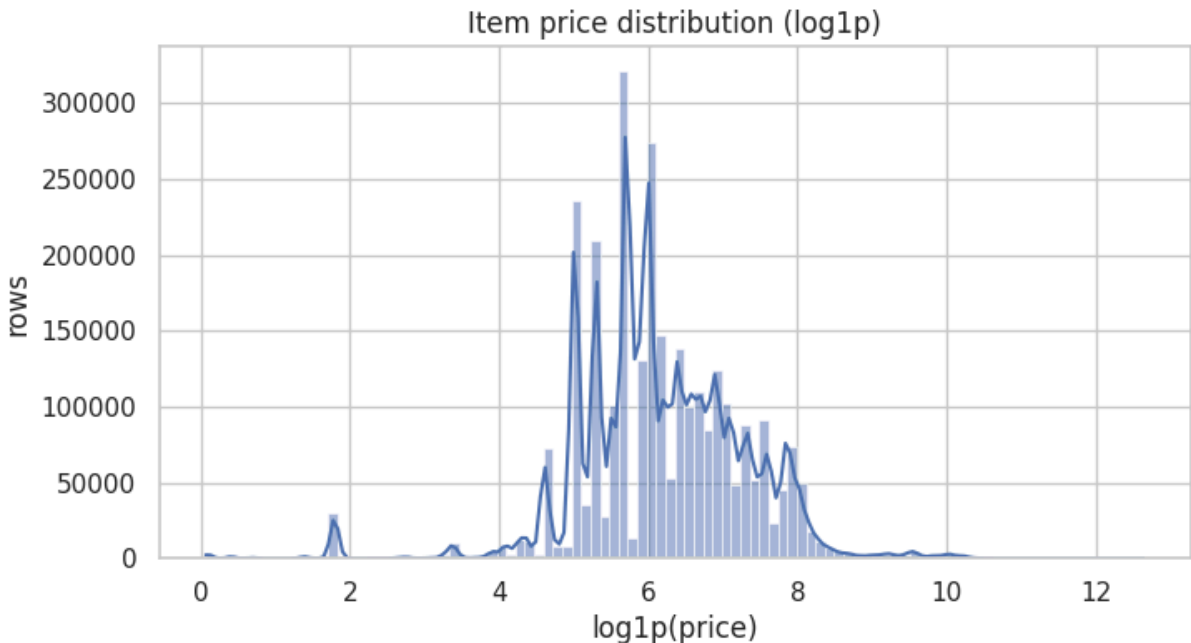
plt.figure(figsize=(8,4))
sns.histplot(np.log1p(df["item_price"]), bins=100, kde=True)
plt.title("Item price distribution (log1p)")
plt.xlabel("log1p(price)"); plt.ylabel("rows")
plt.show()


```

```
# Sparsity: how many days does each (shop,item) pair appear?
pair_days = (df.groupby(["shop_id","item_id"])["date"]
              .nunique()
              .rename("n_days")
              .reset_index())
plt.figure(figsize=(8,4))
sns.histplot(pair_days["n_days"], bins=50)
plt.title("Days available per (shop,item) pair")
plt.xlabel("distinct days"); plt.ylabel("pairs")
plt.show()
```



```
/home/btheard/retail-alpha-forecaster/.venv/lib/python3.10/site-packages/pandas/core/arraylike.py:399: RuntimeWarning: divide by zero encountered in log
lp
result = getattr(ufunc, method)(*inputs, **kwargs)
```







```
In [9]: # --- Cell 5: time signals ---

# Aggregate to daily total units
daily = df.groupby("date", as_index=False)["item_cnt_day"].sum()

plt.figure(figsize=(12,4))
sns.lineplot(data=daily, x="date", y="item_cnt_day")
plt.title("Total units sold per day")
plt.xlabel("date"); plt.ylabel("units"); plt.tight_layout()
plt.show()

# Day-of-week effect (only meaningful with full data window)
df["dow"] = df["date"].dt.day_name()
dow_avg = (df.groupby("dow", as_index=False)["item_cnt_day"]
           .mean()
           .assign(dow=lambda x: pd.Categorical(x["dow"],
           ["Monday", "Tuesday", "Wednesday", "Thursday", "Friday", "Saturday", "Sunday"],
           ordered=True))
           .sort_values("dow"))

plt.figure(figsize=(8,4))
sns.barplot(data=dow_avg, x="dow", y="item_cnt_day")
plt.title("Average units per day by day-of-week")
plt.xlabel("day"); plt.ylabel("avg units"); plt.xticks(rotation=30); plt.tight_layout()
plt.show()
```




```
In [13]: # --- Cell 6: month x shop heatmap (sparsity pattern) ---

month_shop = (df.groupby(["date_block_num", "shop_id"])["item_cnt_day"]
               .sum().reset_index())
pivot = month_shop.pivot(index="shop_id", columns="date_block_num", values="

plt.figure(figsize=(12,6))
sns.heatmap(np.log1p(pivot), cmap="viridis") # log1p to compress scale
plt.title("log1p(units) by shop (rows) and month (cols)")
plt.xlabel("date_block_num"); plt.ylabel("shop_id"); plt.tight_layout()
plt.show()

/home/btheard/retail-alpha-forecaster/.venv/lib/python3.10/site-packages/pa
das/core/internals/blocks.py:393: RuntimeWarning: divide by zero encountered
in log1p
    result = func(self.values, **kwargs)
```



## Executive Summary – Notebook 1 (Exploratory Data Analysis)

In this notebook, we explored the raw sales dataset to understand its structure, quality, and potential for forecasting.

### Key Findings

- **Coverage & Scope**
  - 2.9M rows spanning **Jan 2013 – Oct 2015** (34 months).
  - **60 shops** and **~22k unique items**.
  - Continuous coverage with **no missing months**.
- **Data Quality**
  - **Returns (item\_cnt\_day < 0):** ~7,356 flagged rows.
  - **Negative/zero prices:** 1 negative and 1 zero price row detected.
  - **Duplicate shop-item-date entries:** 28 duplicates found.
  - Overall dataset is clean, with few anomalies.
- **Price Distribution**
  - Most prices fall between **100–5,000**.
  - Extreme outliers exist (100k+), suggesting the need for clipping.
  - Log-transformed prices ( `log1p` ) follow a near-normal distribution.
- **Sales Patterns**
  - **Daily sales:** clear weekly effects and seasonal spikes (late 2014, early 2015).



- **Monthly shop heatmap:** confirms activity gaps and variability across shops.
- **Sparsity:** many shop-item pairs have very short or intermittent sales histories.

## Takeaways for Modeling

- **Forecasting should be monthly**, not daily, due to sparsity.
- Need to **clean anomalies** (returns, duplicates, outlier prices).
- Strong **seasonal and shop-level variability** → valuable signals for features.
- Provides a blueprint for feature engineering in **Notebook 2**.

In [ ]: